

# Conversation between a Julian and non-Julians

2.9.2024 OpenMendel meeting

## Some personal encounters

“How do I use your [Julia] software in R/Python?”

[My collaborator have some C++ code (wrapped by R/Rcpp) that I want to use.] “How do I call it from Julia?”

“Most of genetics and bioinformatics rely on binaries. Can you get a binary from Julia?”

# Today's agenda

I want to share my experience in doing the following

- Writing Python/R package that internally calls Julia code
- Writing C++ wrapper to be called by Julia
- Creating “binary executables” from Julia package

# Today's agenda

I want to share my experience in doing the following

- Writing Python/R package that internally calls Julia code

*This is easy (2-3 days of effort)*

- Writing C++ wrapper to be called by Julia

*This is hard (1-2 month effort)* unless you are C++ expert

- Creating “binary executables” from Julia package

*This is easy (1-2 days effort)* if you are willing to accept >100MB “binaries”

# Write a python/R wrapper for your Julia package

- The packages JuliaCall (juliacall) calls Julia from R (python)
- Our strategy: write a thin python/R package that:
  1. Installs Julia (using JuliaCall)
  2. Tell Julia to install your Julia package
  3. Call into Julia code from R/Python via JuliaCall
- This strategy is used by SciML for DifferentialEquations.jl, I copied them for the Knockoffs.jl package
- For a regular Python/R user, the fact that most code is written in Julia becomes an implementation detail
- Effort required: writing the wrapper is rather quick, but re-writing documentations/unit tests/CI are time consuming. You need documentation to put an R package on CRAN (but PyPI will take anything).

## Example: knockoffsr and knockoffspy for Knockoffs.jl

- These are thin wrapper packages with only 62 (91) lines of source code
- They are regular python/R packages and can be installed in the regular way
  - In python: `pip install knockoffspy`
  - In R:

```
library(devtools)
install_github("biona001/knockoffsr")
```
- Look at diffeqr and diffeqpy for full examples setting up unit test/CI/docs

# Example Usage

- In R

```
> library(knockoffsr)
> ko <- knockoffsr::knockoff_setup()
> result <- ko$modelX_gaussian_group_knockoffs(
  X, "maxent") # assumes X is available
> S <- result$S
```

- In python

```
>>> from knockoffspy import ko
>>> ko.install()
>>> result = ko.modelX_gaussian_group_knockoffs(
  X, "maxent") # assumes X available
>>> S = result.S
```

```
struct GaussianGroupKnockoff{T}
  X::Matrix{T}
  Xko::Matrix{T}
  groups::Vector{Int}
  S::Matrix{T}
  gammas::Vector{T}
  m::Int
  Sigma::Matrix{T}
  method::Symbol
  obj::T
end
```

In Julia, the function  
modelX\_gaussian\_group\_knockoffs()  
returns a struct

## Ben's recommendations for writing Julia code (to make Python/R wrapping easier)

- Use regular English letters for structs, i.e. no Greek letters, latex annotations, or emojis (you cannot type them in R)
- Functions should accept strings rather than symbols (no “symbol” in R/python)
- Functions should not require Julia-specific types, e.g. `::Symmetric` would fail. Rather, it should accept `AbstractMatrix` and you write a separate routine to check symmetry, e.g. `isapprox(X, X', atol=1e-8)`

More than 90% people use R/python. If it is easy to provide an R/python wrapper, why not?



# Writing Julia wrapper for C++ code

- Overall strategy (red = step that was hard for me):
  1. Clone target repository containing C++ code
  2. Write wrapper code (**in C++**) via CxxWrap.jl, upload code to forked repo
    - We will compile it into a shared library (.so) file
  3. Use BinaryBuilder.jl to compile wrapper code into a jll package
    - Compilation for different OS and CPU architectures is possible (**the build script is written by you!**)
  4. Deploy jll package to your personal github, then to Yggdrasil when it is fully ready
  5. To expose functionality, write regular Julia package that calls jll package internally
- In this setup,
  - Users run Julia code the standard Julian way. The fact that your package has C++ dependency becomes an implementation detail
  - Users never compile anything: compiled library (i.e. provided by your jll package) is **downloaded** as an artifact
    - In contrast, R/RCpp make users run a build script to compile the C++ code during package installation, and this step fails very often
  - Extremely small overhead calling C++ code (not humanly noticeable)

# A Full Example

I am providing everything here, even though they are surely not good example, because I had trouble finding a **full** example online

- Basic C++ wrapper
  - [https://github.com/biona001/ghostbasil/blob/master/julia/ghostbasil\\_wrap.cpp](https://github.com/biona001/ghostbasil/blob/master/julia/ghostbasil_wrap.cpp)
- Basic build\_tarballs.jl (from BinaryBuilder.jl):
  - [https://github.com/biona001/Ghostbasil.jl/blob/main/src/build\\_tarballs.jl](https://github.com/biona001/Ghostbasil.jl/blob/main/src/build_tarballs.jl)
  - Run ``julia build_tarballs.jl --deploy="GithubUSERID/SITE_jll.jl"`` to get jll package
- Basic jll package
  - [https://github.com/biona001/ghostbasil\\_jll.jl](https://github.com/biona001/ghostbasil_jll.jl)
- Basic upper level package:
  - <https://github.com/biona001/Ghostbasil.jl>

Everything + explanations:

[https://github.com/biona001/Ghostbasil.jl/blob/main/test/400phenotypes\\_binary\\_executable.ipynb](https://github.com/biona001/Ghostbasil.jl/blob/main/test/400phenotypes_binary_executable.ipynb)

# Ben's recommendation for writing C++ wrappers

- Finding a full example is the most important imo. Somehow, this was difficult
  - If a python/R wrapper exist for the C++ code, study them carefully
- Use CxxWrap.jl instead of Cxx.jl
  - Bart Janssens (@barche) did many helpful CxxWrap.jl tutorial/workshop at JuliaCon. The actual documentation of CxxWrap.jl was not very helpful for me
  - jluna.jl seems promising and well maintained but I have not tried it
- wrapit is C++ package which automatically generates C++ wrapper for Julia, based on CxxWrap.jl. You should try it first, but it never worked for me.
- Specifying compat entry in build\_tarballs.jl for libcxxwrap\_julia\_jll is extremely important, e.g.

```
Dependency("libcxxwrap_julia_jll"; compat = "0.11.2")
```

CxxWrap.jl internally depends on this binary, but it **often** release breaking changes

- The step calling BinaryBuilder.jl can be complicated. Fortunately, thousands of example `build\_tarballs.jl` files can be found at Yggdrasil. Unfortunately, most of them copy each other.

Many R/Python packages are internally calling C++ code, but as of 2024 it is not easy to write Julia wrappers for them, let alone automate it

# Building binaries from Julia package

- Many users prefer binaries:
  - No required installation/compilation (just download a pre-compiled program)
  - Clearly defined inputs/outputs
  - E.g. PLINK, bcftools, ADMIXTURE, ... etc
- Strategy: use `create\_app` function of PackageCompiler.jl
  - StaticCompiler.jl is another option, but it seems less flexible (I have not tried it)
- PackageCompiler.jl produces **a folder** which includes the desired executable + all the necessary files to run it
  - When zipped, it is 100-400 MB (seems like this number can be optimized)
- The entire folder can be sent to another machine directly, e.g. one without Julia installed
  - The target machine **must have the same operating system**
  - Unlike BinaryBuilder.jl, one needs to compile for Mac, Linux, Windows separately (e.g. via AWS?)

# Strategy

1. Write a regular Julia package
2. Define a function `julia_main()` which operates on `ARGS` and returns
  - 0 (if program successfully ran) or
  - 1 (if an error occurred)
3. Call `PackageCompiler.create_app``
  - `src` = directory of your package
  - `des` = output folder
  - The result is an executable whose usage is `PKG_NAME ARGS[1] ARGS[2] ...`
  - `include_lazy_artifacts` keyword is required if your package has external (artifact) dependencies, e.g. one based on a jll package

```
function julia_main()::Cint
    try
        # do something based on ARGS?
    catch
        Base.invoke_latest(Base.display_error, Base.catch_stack())
        return 1
    end
    return 0 # if things finished successfully
end
```

```
using PackageCompiler, GhostKnockoffGWAS
src = normpath(pathof(GhostKnockoffGWAS), "../..")
des = normpath(pathof(GhostKnockoffGWAS), "../..linux_64")
@time create_app(src, des,
    include_lazy_artifacts=true,
    force=true,
)
```

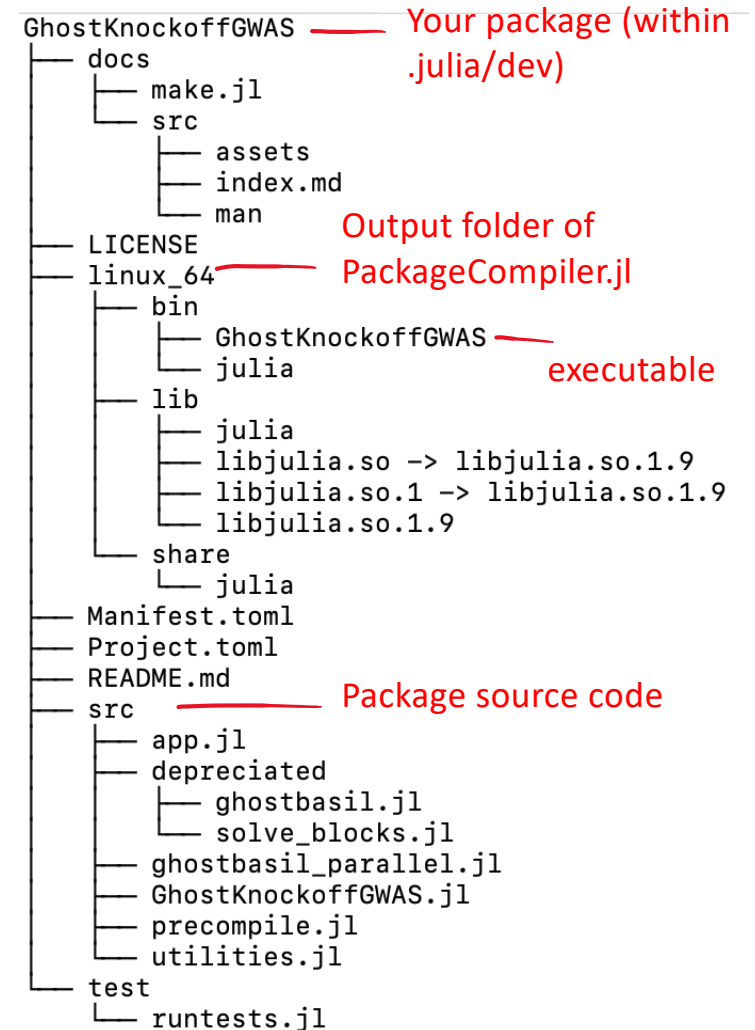
# Example

- `create\_app()` will run for 1-2h and then generate a folder containing the executable + bundle of other files → →
- The executable can be invoked in the terminal directly, e.g.

```
$ GhostKnockoffGWAS example_zfile.txt EUR 506200 38 example_output
```

- For more command-line-program-like usage, use ArgParse.jl to parse `ARGS`:

```
$ GhostKnockoffGWAS \  
  --zfile example_zfile.txt \  
  --LD-files EUR \  
  --N 506200 \  
  --genome-build 38 \  
  --out example_output
```



# Ben's recommendation for using PackageCompiler.jl

- Compiling Julia code into an app (binary? Executable?) is pretty easy
  - The result is not a “true binary”, whatever that means
- Beware of accidentally depending on large binaries (e.g. MKL\_jll.jl is 600+MB)
- Underlying Julia code could still be **JIT compiled**, if you are not careful with precompilation statements
  - E.g. My “time to help statement” now takes 0.5 sec (still slow) but an improvement from 10s

```
$ time GhostKnockoffGWAS -h
real    0m0.574s
user    0m0.343s
sys     0m0.159s
```

- Still not easy to compile for every platform (e.g. Windows/Mac/linux), not sure why it can't work like BinaryBuilder.jl
- Still not clear how to distribute the resulting (large) binaries

# Summary

- Writing Python/R package that internally calls Julia code?

This is easy (2-3 days of effort), but re-writing documentation/tests/CI is a lot more effort

- Writing C++ wrapper to be called by Julia?

This is hard (1-2 month effort) unless you are C++ expert. It is probably easier if you re-implement everything in Julia

- Creating “binary executables” from Julia package?

This is easy (1-2 days effort) if you are willing to accept >100MB “binaries”. You might encounter small inconveniences but I’m satisfied overall